

Governed AI Policy Response Assistant

Independent portfolio implementation by Shalinee Singh | Microsoft Copilot Studio Workflow | Evidence-based status

BUSINESS PROBLEM	IMPLEMENTED PROTOTYPE
Policy-response drafting is slow and inconsistent; fully autonomous customer communication creates accuracy, privacy, legal, and reputational risk.	Manual inquiry inputs feed a governed Agent, then Human Review, an approval condition, and separate approved/revision outcome records.

Workflow

Manual input	Governed Agent	Human Review	Approval decision	Outcome record
--------------	----------------	--------------	-------------------	----------------

Governance controls

Untrusted input	Customer inquiry cannot override the Agent role or safeguards.
Grounded context	Only three embedded synthetic policies may support the draft.
High-risk escalation	Legal, privacy, fraud, security, discrimination, and regulatory matters require review.
No autonomous send	The prototype contains no direct external-email action.
Accountability	Approval decision and reviewer comments feed separate outcome records.

Validation status

Published workflow	PASSED	Workflow health check and publication completed.
Agent node test	PASSED	Returned a safe identity-policy draft, policy basis, and review note.
Outlook review request	DELIVERED	Interactive Yes/No and reviewer-comment form rendered.
Approval callback	BLOCKED	Tenant endpoint returned HTTP 400.
Teams notification	BLOCKED	Managed service returned HumanInTheLoopNotificationFailed.
Complete workflow run	BLOCKED	Environment had no available Copilot Credits.

AI transformation contribution

Use-case discovery; requirements and acceptance criteria; process redesign; Copilot Studio configuration; prompt and instruction design; responsible-AI controls; human-review workflow; approval branching; risk register and RACI; UAT planning; evidence management; and diagnosis of licensing, capacity, connector, and tenant-policy constraints.

Production roadmap

Obtain approved capacity; move policies to a managed knowledge source; validate review callbacks; add correlation IDs, durable audit records, monitoring, timeouts, exception handling, solution packaging, environment promotion, and full UAT; then obtain business, legal, privacy, security, records, architecture, and platform approvals.